

Using Voice and Biofeedback to Predict User Engagement during Requirements Interviews

Alessio Ferrari, Thaide Huichapa, Paola Spoletini, Nicole Novielli, Davide Fucci, and Daniela Girardi

Abstract—Capturing users' engagement is crucial for gathering feedback about the features of a software product. In a market-driven context, current approaches to collect and analyze users' feedback are based on techniques leveraging information extracted from product reviews and social media. These approaches are hardly applicable in bespoke software development, or in contexts in which one needs to gather information from specific users. In such cases, companies need to resort to face-to-face interviews to get feedback on their products. In this paper, we propose to utilize biometric data, in terms of physiological and voice features, to complement interviews with information about the engagement of the user on product-relevant topics. We evaluate our approach by interviewing users while gathering their physiological data (i.e., *biofeedback*) using an Empatica E4 wristband, and capturing their voice through the default audio-recorder of a common laptop. Our results show that we can predict users' engagement by training supervised machine learning algorithms on biometric data ($F1=0.72$), and that voice features alone are sufficiently effective ($F1=0.71$). This work is one of the first studies in requirements engineering in which biometrics are used to identify emotions. Furthermore, this is one of the first studies in software engineering that considers voice analysis. The usage of voice features can be particularly helpful for emotion-aware requirements elicitation in remote communication, either performed by human analysts or voice-based chatbots, and can also be exploited to support the analysis of meetings in software engineering research.



1 INTRODUCTION

The development of novel software products, as well as the improvement of existing ones, can deeply benefit from the involvement of users in requirements engineering (RE) activities [1]. Getting feedback from the user base has been recognised to lead to increased usability, improved satisfaction [2], better understanding of requirements [3], and creation of long-term relationships with customers [4].

User feedback can take implicit and explicit forms, and different means are available to collect this information. In particular, data analytics applied to user opinions and to usage data has seen an increasing interest in the last years, leading to the birth of RE sub-fields such as *crowd RE* [5], [6] and *data-driven RE* [7], [8]. In the case of bespoke development (i.e., when customer- or domain-specific products' requirements need to be engineered), it is still common to resort to traditional RE practices, such as prototyping, observations, usability testing, and focus groups [9]. Among the classical techniques, user interviews are one of the most commonly used to gather requirements and feedback [10], [11], [12]. A recent study involving 12 software develop-

ment companies showed that individual interviews are the second most used elicitation technique, preceded by group workshops and focus groups [13]. Several aspects have been observed to influence the success and failure of interviews, such as the domain knowledge of the requirements analyst [12], [14], ambiguity in communication [15], and typical mistakes such as not providing a wrap-up summary at the end of the interview session, or not creating rapport with the interviewee [16].

Currently, little attention is dedicated to the emotional aspects of interviews and, in particular, to users' *engagement*. Capturing engagement is crucial for gathering feedback about the features of a certain product, and have a better understanding of the users' preferences. The field of *affective RE*—which combines ideas from Affective Computing with RE practices—recognised the role of users' emotions and studied it extensively. Contributions include applications of sentiment analysis to app reviews [17], [18], analysis of users' facial expressions [19], [20], the study of physiological reactions to ambiguity [21], and the augmentation of goal models with user emotions elicited through psychometric surveys [22].

In this paper, we aim to extend the body of knowledge in affective RE by studying users' emotions during interviews. We focus on *engagement*—i.e., the degree of positive or negative interest on a certain product-related aspect discussed in the interview. We perform a study with 31 participants taking part in a simulated interview during which we capture their biofeedback using an Empatica E4 wristband, we record their voice through a common laptop recorder, and collect their self-assessed engagement. We compare different machine learning algorithms to predict users' engagement based on features extracted from biofeedback and voice

- Alessio Ferrari is with the Institute of Information Science and Technologies (ISTI) of the Italian National Research Council (CNR), Pisa, Italy. E-mail: alessio.ferrari@isti.cnr.it
- Thaide Huichapa and Paola Spoletini are with Kennesaw State University (KSU), GA, USA. E-mail: thuichap@students.kennesaw.edu, pspoleti@kennesaw.edu
- Nicole Novielli and Daniela Girardi are with the University of Bari, Bari, Italy. E-mail: {nicole.novielli, daniela.girardi}@uniba.it
- Davide Fucci is with Blekinge Tekniska Högskola (BTH), Karlskrona, Sweden. E-mail: davide.fucci@bth.se

Manuscript received 10 Sept., 2021; revised ...

signals.

Our experiments show that engagement can be predicted in terms of valence and arousal [23] with an F1-measure of, respectively, 63% and 65%, when only considering biofeedback signals. When using voice signals alone, the performance in terms of F1-measure increases to 70% and 71%, showing that voice features can be strongly predictive of users' engagement. The combination of biofeedback and voice features leads to the best performance for both valence and arousal—F1-measure of 72%.

This paper builds upon a previous conference contribution by the same authors [24], in which only the biofeedback signals were used for prediction. The current paper replicates and expands the experiments. In particular, we introduce additional biometric features, based on voice signals, as well as additional data preparation options—namely standard scaling, oversampling and data imputation—which allow us to improve the previous results, even using only voice features. This is a particularly relevant outcome, as voice analysis can be scaled, with minimal costs, to a large number of interviewees, who may be located remotely. The idea could be particularly effective in case voice-based chatbots are used as interviewers, to equip them with the ability to detect the engagement of the user from their emotional prosody [25], and steer the interview accordingly.

This paper contributes to the literature with empirical work in which we show that biometric features, including physiology- and voice-related metrics, can be applied to predict users' engagement during requirements interviews. While biofeedback analysis has been previously applied in other software engineering areas [26], [27], [28], this is one of the first works that explores its potential in RE. Furthermore, to our knowledge, this is the first work that uses voice analysis in software engineering, thus opening to studies focused on speech-intensive activities of software development. A replication package is also made available¹ [29] to enable other researchers to build on our results.

The remainder of the paper is structured as follows. In Section 2, we present background definitions of engagement and theories of emotions, as well as related work in RE and software engineering. In Section 3, we report our study design, whereas Section 4 reports its results. We discuss the implications of our study in Section 5 and its limitations in Section 6. Finally, Section 7 concludes the paper.

2 BACKGROUND AND RELATED WORK

In this section, we first clarify the relationship between emotion modelling and engagement (Sect. 2.1). Then, we present the background on affect modelling and emotion classification using biofeedback (Sect. 2.2) and voice analysis (Sect. 2.3). Finally, we discuss relevant related work in RE and software engineering (Sect. 2.4).

2.1 Engagement and Emotions

Affective states vary in their degree of stability, ranging from personality traits—i.e., long-standing, organized sets of characteristics of a person—to emotions—i.e., transient

and typically episodic, dynamic, and structured events. Emotions involve perceptions, thoughts, feelings, bodily changes, and personal dispositions to experience them. Emotions are episodic and dynamic in that, over time, they can vary depending on several factors [30].

Psychologists have investigated the nature and triggers of emotions for decades, leading to a plethora of theories of emotions. Among these theories, cognitive models describe emotions as reactions to cognition. For example, the OCC model [31] defines a taxonomy of emotions and identifies them as *valenced* reactions (either positive or negative) to the cognitive processes involved in evaluating objects, events, and agents. Analogously, Lazarus describes nine negative (Anger, Anxiety, Guilt, Shame, Sadness, Envy, Jealousy, and Disgust) and six positive (Joy, Pride, Love, Relief, Hope, and Compassion) emotions, as well as their *appraisal* patterns: when a situation is congruent with the person's goals positive emotions arise; otherwise, negative emotions are triggered when one's goal are threatened [32].

In line with these theories and consistently with the operationalization adopted in our previous study [24], we use emotions as a proxy for users' *engagement* during interviews. When evaluating the importance of a topic during an interview, the appraisal process of an individual is responsible for triggering an emotional reaction based on the perceived relevance of the topic with respect to his/her goal, values, and desires. Our choice is further supported by previous empirical findings demonstrating how emotions can be leveraged for detecting engagement in speech-based analysis of conversations [33] or to detect students' motivation [34].

Consistently with prior research on emotion awareness in software engineering [26], [35], [36], we adopt a dimensional representation of developers' emotions. In particular, we refer to the Circumplex Model of Affect by Russell [23], which models emotions along two continuous dimensions, namely *valence*, that is the pleasantness of the emotion stimulus, ranging from pleasant to unpleasant, and *arousal*, that is the level of emotional activation, ranging from activation to deactivation. Pleasant emotional states, such as happiness, are associated with *positive* valence, while unpleasant ones, such as sadness, are associated with *negative* valence. The arousal dimension captures, instead, the level of emotional activation. Some emotions are associated with the person being inactive, thus experiencing *low* arousal, as in *sadness* or *relaxation*. Conversely, high level of arousal are associated to high emotional activation, as in *anger* or *excitement*.

We expect to observe different forms of engagement in relation to valence and arousal: positive-high engagement (i.e., positive valence and high arousal) may occur when users discuss topics that they consider relevant and towards which they have a positive feeling, e.g., a feature users like and have an opinion they want to discuss about; negative-high engagement (i.e., negative valence and high arousal) may occur when topics are relevant but more controversial, such as a feature that users do not like, or a bug they find annoying. Low engagement may occur when the user does not have a strong opinion on the topic of the discussion, and is either calm (positive valence, low arousal) or bored by the conversation (negative valence, low arousal).

1. <https://github.com/alessioferrari/VoiceBiofeedEmo>

2.2 Biofeedback-based Classification of Emotions

The use of physiological signals for emotion recognition has been largely investigated by affective computing research [37], [38], [39], [40], [41]. Previous work studied the relationship between emotions and biometrics such as the electrical activity of the brain—e.g., using electroencephalogram (EEG) [40], [42], [43], [44], the electrical activity of the skin, or electrodermal activity (EDA) [45], [46], the electrical activity of contracting muscles measured using electromyogram (EMG) [39], [41], [47], and the blood volume pulse (BVP) from which heart rate (HR) and its variability (HRV) are derived [37], [48].

Electroencephalogram (EEG) captures the electrical activity of the brain through electrodes placed on the surface of the scalp. Changes in the EEG spectrum correlate with increased or decreased overall levels of arousal or alertness [42] as well as with the valence of the emotion experienced [40], [43].

The electrodermal activity (EDA) measures the electrical conductance of the skin due to the sweat glands activity. EDA correlates with the arousal dimension [49] and its variation occur in presence of emotional arousal and cognitive workload. Hence, EDA has been employed to detect excitement, stress, interest, attention as well as anxiety and frustration [45], [46].

Heart-related metrics have been successfully employed for emotion detection [37], [48]. In particular, blood volume pressure (BVP) measures the changes in the volume of blood in vessels, while Heart Rate (HR) and its Variability (HRV) capture the rate of heart beats. Significant changes in the BVP are observed in presence of increased cognitive and mental load [50]. Increases in HR occur when the body needs a higher blood supply, for example in presence of mental or physical stressors [51].

Finally, several studies demonstrated the high predictive power of facial EMG for emotion recognition [39], [47]. However, it leads to poor results when the sensors are placed on body parts other than the face (i.e., the arms [41]).

In a recent study, Girardi et al. [26] identify a minimum set of sensors including EDA, BVP, and HR for valence and arousal classification. To collect such physiological signals, they use the Empatica E4 wristband and detect developers' emotions during software development tasks. They found that the performance obtained using only the wristband are comparable to the one obtained using an EEG helmet together with the wristband.

Accordingly, in this study we use EDA, BVP, and HR collected using Empatica E4, a noninvasive device that participants can comfortably wear during interviews (see Section 3.2), thus increasing the ecological validity of our study. Furthermore, we combine biofeedback with *voice* features, which were not considered in previous works.

2.3 Voice Analysis and Classification of Emotions

Classification of emotions based on the analysis of voice features is a well-developed research field, normally referred as *speech emotion recognition* (SER). Different surveys have been recently published on the topic [52], [53], [54], which highlight the maturity of research, but also point out the limits in terms of real-world applications, mainly due to

limited gold standard datasets available for SER systems' training and assessment.

Speech is composed by a diverse set of acoustic features, and its information content is usually accompanied by other so-called supporting modalities, including linguistic features (i.e., the textual content equivalent to a verbal utterance), visual features, and physiological signals such as those discussed in the previous section.

Acoustic features used in SER are normally classified into *prosodic* (e.g., pitch, tone), *spectral* (i.e., frequency-based representation of the sound produced), *voice quality* (e.g., measuring the stability of the voice) and Teager energy operator (TEO)-based features, specifically developed to detect stress from the voice signal. Prosodic and spectral features are the most commonly used in the literature [52], [54]. In particular, most commonly used features are Mel-scaled spectrogram and Mel-frequency cepstral coefficients (MFCCs), which are spectral features that mimic the reception pattern of sound frequency intrinsic to a human [55]. [55] uses also Chromagrams—typically used for music representation—since the other features were recognised to be poor in distinguishing pitch classes and harmony [56].

Research in SER initially focused on identifying relevant features and combination thereof to optimise the performance of traditional classification algorithms [57], [58], leading to good recognition rates especially with Support Vector Machines [59], [60]. With the advent of deep neural networks, and the possibility of overcoming the feature engineering problem altogether, the focus shifts to the selection of appropriate network architectures, and promising results are achieved through Convolutional Neural Networks (CNNs) and Long Short-Term Memory (LSTM) models [61], [62]. To address the problem of data scarcity, the recent development of transfer learning methods have been also experimented [63]. Another avenue of research in SER, still under development, is the combination of pure audio features with other contextual cues, including videos [63], [64], [65].

Applications of the SER techniques are mostly in the field of human-computer interaction (HCI). Ramakrishnan and El Emery [66] lists a set of ten different possible applications, including lie detection, treatment of language disorders, driving support systems [60], surveillance [67], and smart assistants. These latter have become commercially available in recent years—e.g., Siri, Cortana—and represent one of the natural area of exploitation of SER research. Another traditional application, closely related to our context, is the recognition of customer's emotions during conversations with call center operators [68], [69], [70], [71], [72]. These works are oriented to identify critical phases of the dialogue between a customer and an artificial operator. This can be generally useful to understand when the artificial agent is irritating the customer, thus to deciding to transfer the call to a human operator.

2.4 Related Work

Biometric sensors have been leveraged in several software engineering studies for recognizing cognitive and affective states of software developers.

In one of the early studies in this field, Parnin [73] envisions an approach to infer developers' cognitive states based

on the analysis of sub-vocal signals, that is the electrical signal the brain transfers to the mouth and vocal cords while performing complex cognitive activities. While presenting preliminary findings, this study demonstrate that it is possible to use EMG to capture the subvocalization associated to programming and leverage this information to distinguish between easy and hard development tasks. Fucci et al. [74] use multiple biometrics, including EEG, EDA, HR, and BVP, to distinguish between code and prose comprehension tasks during software development. Fritz et al. [27] employ EEG, BVP, and eye tracker to measure the difficulty of programming tasks, thus preventing developers from introducing bugs. The authors use the same set of sensors in a follow-up work aimed at classifying the emotional valence of developers during programming tasks [35]. Girardi et al. [26] replicate previous findings by Müller and Fritz [35] regarding the use of non-invasive sensors for valence classification during software development tasks. Furthermore, Girardi et al. also address the classification of arousal. In a recently published work, Girardi et al. [28] study the relationship between biometric signals and productivity of developers with an *in vivo* setting involving engineers at work. Other studies also propose approaches for predicting developers' interruptibility [75] and for identifying code quality concerns [76] by leveraging EDA, HR, and HRV.

Biofeedback has been used also in RE, mainly to capture users' emotions *while* using an app. For example, some studies [20], [77] use mobile phone cameras to recognize facial muscle movements and associate them to the users' emotions when using different features of an app. This methodology was recently applied to enable user validation of new requirements [19] and to for identifying usability issues [78] with minimal privacy concerns [79]. Part of the authors of the current paper previously proposed using biometrics in requirements elicitation interviews [21]. Our previous work focused on ambiguity, and remained at the research preview stage, as it evolved in the current work after pilot experiments. Other work in affective RE acknowledge the primary relevance of users' emotions [80], and generally focus on the application of sentiment analysis techniques to textual user feedback, for example from tweets or app reviews [81], [82]. Other uses of sentiment analysis in RE include the prediction of tickets escalation in customer support systems [83]. Emotions have been also considered in early-stage RE activities, such as requirements elicitation and modelling. For example, Colomo-Palacios et al. [84] asked users to rank requirements according to Russel's Valence-Arousal theory, which is the one that we adopt in the present study. Other researchers leverage information regarding users' emotions gathered through psychometrics (e.g., surveys) to augment traditional requirements goal modelling approaches [22], [85] and artefacts, such as user stories [86].

2.4.1 Contribution

Compared to previous works using biofeedback and voice analysis in software engineering and RE, our study is among the first ones to specifically focus on users' interviews rather than product usage or development tasks. Previous studies (e.g., [20], [77]) focus on detecting the user's engagement experienced *while* using the software features. In our case,

we aim to detect users' engagement about certain features when users reflect on the features and speak about them. This captures a different moment—a verbalized, more rational one—of the relationship between the user and the product. Furthermore, in interviews we can consider *what if* scenarios (e.g., financial and privacy-related questions in Table 1), which is not possible when performing observations without interacting with users. Finally, to our knowledge, our work is among the first ones that use voice features to predict the emotion of a speaker involved in a software engineering activity.

3 RESEARCH DESIGN

Our study is *exploratory* in nature, aiming to investigate a certain area of interest—i.e., *engagement* in interviews—and identify possible avenues of research. We adopt a quantitative experimental approach involving human subjects, and oriented to compare software-based artifacts (i.e., machine learning algorithms and feature configurations). The study was approved by the Kennesaw State University review board (study 16-068).

The main goal of this study is to understand to what extent we can use biofeedback devices and voice analysis to predict users' engagement during interviews. Accordingly, we formulate the following research questions (RQs).

- **RQ1:** *To what extent can we predict users' engagement using biofeedback measurements and supervised classifiers?* With this question, we aim to understand whether it is possible to automatically recognize engagement with biofeedback. More specifically, we want to assess to what extent we can recognize emotional valence and arousal—i.e., the two dimensions we use for the operationalization of engagement. To collect training and testing data, we first interview Facebook users², asking their opinion about the platform. After the interview, we ask them to report their engagement for each of the different questions. During the interviews with users, we acquire their raw biofeedback signals. We use features extracted from the signals, and consider intervals of reported engagement as classes to be predicted. Based on these data, we evaluate and compare different supervised machine learning classifiers.
- **RQ2:** *To what extent can we predict users' engagement using voice analysis and supervised classifiers?* This question aims to understand whether we can recognize engagement with automatic voice analysis. This RQ is motivated by the need to support a different use case in which biofeedback devices might not be used. It is the case, for example, of interviewees being uncomfortable wearing a tracking device during the conversation. In other situations the use of physical devices might even be unfeasible, due to remote interview settings. To this extent, we record the audio of the interviews, and we extract voice features from the audio signals. We then use the voice features to

2. Although our study is not primarily oriented to consumers' products, selecting Facebook as discussion topic facilitates the selection of participants.

train and compare the previously used supervised classifiers.

- **RQ3:** *To what extent can we predict users' engagement by combining voice and biofeedback features?* This question aims to investigate the use of voice and biofeedback features in conjunction. By training and comparing the classifiers, we check if and in which way the combination of features improves the performance of the approach.

3.1 Study Participants

We recruited 31 participants among the students of Kenesaw State University with an opportunistic sampling. The participation was not restricted by major or academic level, but the only main requirement was to be an active Facebook user (access to Facebook at least once per day, self-declared), as the user interview questions dealt with this social network. More than 90% of the participants were undergraduate students divided in 11 majors. To account for differences in biometrics due to physiological aspects [87], we try to have a pool of participants as varied as possible by including multiple ethnic groups and both female and male subjects. Specifically, approximately 65% of the participants were male, and their age varied between 18 and 34 with both median and average equal to 22. Participants were either native speakers or proficient in English. The majority (58%) were white/Caucasian, 23% black/African American, 13% Hispanic/Latino, and the remaining 6% was Asian/Pacific islander. During the data analysis, we removed 10 participants because either the collected data were incomplete or the available information were not considered reliable (e.g., they provided the same response to all the questions in the surveys). Of the remaining 21 participants, approximately 67% were male with the following racial/ethnicity distribution, 67% white/Caucasian, 14% black/African American, 14% Hispanic/Latino, and 5% Asian/Pacific islander. We collected information about the ethnicity of participants because the research demonstrated that heart-rate optical sensors might give more/less reliable readings based on the skin tones. Having a diverse pool of participants in terms of ethnicity strengthens the validity of our empirical findings. Participants received a monetary incentive of \$25 for up of one hour of their time.

3.2 Biofeedback Device and Signals

The device we use to acquire the biofeedback is the Empatica E4³ wristband. We selected it as it is used in several studies in affective computing [51] as well as in the field of software engineering [35], [74]. Using the Empatica E4, we collected the following signals:

- **Electrodermal Activity:** EDA can be evaluated based on measures of skin resistance. Empatica E4 achieves this by passing a small amount of current between two electrodes in contact with the skin, and measuring electrical conductance (inverse of resistance) across the skin. EDA is considered a biomarker of individual characteristics of emotional responsiveness

and, in particular, it tends to vary based on attentive, affective, and emotional processes [88].

- **Blood Volume Pulse:** BVP is measured by Empatica E4 through a photoplethysmography (PPG)—an optical sensor that senses changes in light absorption density of the skin and tissue when illuminated with green and red lights [89], [90].
- **Heart Rate:** HR is measured by Empatica E4 based on elaboration of the BVP signal with a proprietary algorithm.

Research identified a minimal set of biometrics for reliable valence and arousal detection, consisting in the EDA, BVP, and HR measured by the E4 wristband [26].

3.3 Audio Device and Signals

The interviews' audios were captured using the default audio recorder of a Mac OS laptop, and the files were stored in the classical Waveform Audio File Format (i.e., .wav), which is an uncompressed representation of the raw signal. Audio is a complex, information-rich signal, and a largely variable set of classical features are used to characterise its salient aspects [91]. Among these features, we consider the following ones:

- **Mel Spectrogram:** the acoustic time-frequency representation of sound.
- **Mel-Frequency Cepstral Coefficients (MFCC):** MFCC are the representation of the short-term power spectrum of sound. More in details, Cepstral features contain information about the rate changes in the different spectrum bands and they have the ability to separate the impact of the vocal cords and the vocal tract in a signal. In the MFCC, these features are extracted at the frequency more audible by human ears.
- **Chromagram:** the Chromagram is a 12-element feature vector indicating how much energy of each pitch class is presented in the signal. This is typically used to model harmonic and melodic characteristics of music, and it is recognised as useful also to model the emotional aspect of voice [91].

We choose these features as they are amongst the most common in speech emotion recognition [55], [55], [91].

3.4 Experimental Protocol and Data Collection

Three main roles are involved in the experiment: *interviewer*, *user*, and *observer*. The interviewer leads the experiment by asking questions to the user, while the observer tracks the interview by annotating timestamps of each question, monitoring the output of the wristband, checking that audio recording is operational, and annotating general observations on the interview and behaviour of the user.

The experimental protocol consists of four phases (i) device calibration and emotion triggering, (ii) user's interview, (iii) self-assessment questionnaire, and (iv) wrap-up. At the beginning of the experiment, we explained the different steps to the participants, who had the opportunity to ask clarification questions throughout the experiment. According to the IRB that approved the experiment protocol, the collected data does not present more than minimal risks to the participants.

3. <https://www.empatica.com/research/e4/>

3.4.1 Device calibration and emotion triggering

In line with previous research [26], [35] we run a preliminary step for device calibration and emotion elicitation. The purpose of this phase is threefold. First, we want to check the correct acquisition of the biofeedback signal by letting the wristband record the raw signals for all sensors under the experimenter scrutiny. Second, the collected data will be needed to adjust the scores obtained during the self-assessment questionnaire (see Sect. 3.5). Third, we want the participants to get acquainted with the emotion self-report task.

Accordingly, we run a short emotion elicitation task using a set of emotion-triggering pictures. Each participant watches a slideshow of 35 pictures. Each picture is displayed for 10 seconds, with intervals of five seconds between them to allow the user to relax. The whole slideshow lasts for nine minutes. During the first and last three minutes, calming pictures are shown to induce a neutral emotional state, while during the central 3 minutes the user sees pictures aimed at triggering negative and positive emotions. The pictures have been selected from the Geneva database [92], previously used in software engineering studies [35]. The user is then asked to fill a form to report the degree of arousal and valence they associated to the pictures on a visual scale from 0 to 100. As done in previous work [35], for each picture, the user is asked two questions: “*You are judging this image as:*” (0 = Very Negative; 50 = Neutral; 100 = Very Positive); “*Confronted with this image you are feeling:*” (0 = Relaxed, 50 = Neutral, 100 = Stimulated).

3.4.2 User’s Interview

A trained interviewer conducted the interview with each participant. The interview script consists of 38 questions concerning the Facebook platform. Questions are grouped into seven topics—i.e., *usage habits*, *privacy*, *procedures*, *relationships*, *information*, *money*, and *ethics*. The questions are reported in Table 1. For each topic, we include multiple questions, to allow users sufficient time to get immersed in the topic, and collect more stable biometric parameters in relation to the topic. Questions related to topics we expect to raise more engagement, (i.e., *privacy*, *relationship*, *money*, and *ethics*) are separated by questions on topics that are expected to reduce user engagement (i.e., *usage habits*, *procedures*, and *information*). The lower degree of engagement for the latter topics was assessed during preliminary experiments in which the questions were drafted and finalised⁴. Overall, our instrument has been designed to provoke mild emotions, triggered by reflections about the product features, so to measure if these product-related emotions could be captured by biometric data. We have chosen to use a structured interview questionnaire to have comparable interviews, so to facilitate subsequent analysis. To improve the degree of realism of the interview, we did not constraint the time slot for answering each question, so that the interviewees could freely express their opinion.

During the interview, the wristband records the biofeedback parameters, the audio recorder acquires the voice of

the speakers, while the observer annotates the timestamp of each question. We use this information to align the sensor data with the questions. Based on a preliminary run, each interview was estimated to last for about 20 minutes.

3.4.3 Self-assessment Questionnaire and Wrap-up

For each question in the interview script (i.e., Q_i), the interviewer asks the participant to report their involvement using two 10-point rating scale items: $q_A(Q_i)$: How much did you feel involved with this topic? (1 = Not at all involved; 10 = Extremely involved); $q_V(Q_i)$: How would you rate the quality of your involvement? (1 = Extremely negative; 10 = Extremely positive). These two questions aim at measuring the engagement of the user in terms of *arousal* (q_A) and *valence* (q_V). Afterwards, the observer downloads and stores the wristband data as well as the voice recording and the questionnaires filled by the participant. The wristband memory is then erased to allow further recording sessions.

3.5 Pre-processing, Feature Extraction, Classification

The data from the interview questionnaire are used to produce the gold standard—i.e., the labels for valence and arousal to be predicted.

We define *positive*, *negative*, and *neutral* labels for valence, and *high*, *low*, and *neutral* labels for arousal. We discretize the scores in the rating scale following an approach utilized in previous research [26], [35]. First, we adjust the valence and arousal scores based on the mean values reported while participants watched emotion-triggering pictures (see Section 3.4.1). This step is necessary to take into account fluctuations due to individual differences in the interpretation of the scales in the interview questionnaire. Then, we perform a discretization of the values into the three categories (i.e., labels) for each dimension using k-means clustering⁵.

To synchronize the measurement of biofeedback and voice signals with the self-assessment, we (1) save the timestamp corresponding to the interviewer asking question Q_i (i.e., $timestamp(Q_i)$), (2) calculate the timestamp associated to the next question Q_{i+1} ($timestamp(Q_{i+1})$), and (3) select each signal samples recorded between $timestamp(Q_i)$ and $timestamp(Q_{i+1})$.

For each interview question Q_i , we have:

- a set of biofeedback signal samples (for EDA, BVP and HR) within the time interval associated to Q_i ;
- a voice signal sample in the form of a .wav file—the segment of the .wav file of the whole interview for the time interval associated to Q_i ;
- two labels, one representing arousal ($q_A(Q_i)$) and the other representing valence ($q_V(Q_i)$) according to the self-assessment questionnaire.

The labels are used to form the gold standard to be predicted by the algorithms based on features extracted from the signal samples.

To maximize the signal information and reduce noise caused by movements, we apply multiple filtering techniques. Regarding BVP, we extract frequency bands using

4. During the experiments reported in this paper, we saw that *usage habits* was associated with higher engagement, instead. Discussion on this aspect is reported in Sect. 5.

5. We use the k-means implementation in by the *arules* R package.

TABLE 1
List of questions asked during the Interview Phase

USAGE HABITS
1. Do you use the Facebook chat function?
2. (If yes to 1) Who are the people you talk to most frequently using the Facebook chat? (If no to 1) Do you use any other chat applications?
3. How many hours do you use Facebook per day?
4. When you check Facebook, what is the average length of time you spend per session?
5. Is Facebook your primary source of social media? (If yes, why? If no, what other social media you use more often? Why is it superior?)
PRIVACY
6. If someone shared a photo of you in an embarrassing, incriminating, or shameful situation, how would you react? (Do you think Facebook has a responsibility to prevent it from happening? Should they be allowed to remove the photo on your behalf?)
7. If someone tagged you in a post which contained topics you are not comfortable sharing on Facebook (e.g., your political view, sexual preference, . . .), how would you react? (Do you think Facebook has a responsibility to prevent it from happening?)
8. How would you feel knowing that someone (e.g. your SO) accessed your profile and searched it?
9. Imagine Facebook begins using profile information to generate ad content. Would you be okay with this? (why?)
10. In relation to Facebook, what is private information?
PROCEDURE
11. Can you explain me how to add a new friend on Facebook?
12. Can you explain me how to find Facebook pages that match your interest?
13. Can you explain to me how to block a person on Facebook?
RELATIONSHIP
14. Are you connected on Facebook with members of your family? (If so, do you interact with them using Facebook? If not, why?)
15. Have you ever had a family member (even of your extended family) delete you from his/her friend list? (If so why?)
16. Have you ever wanted to delete or deleted a family member (even of the extended family) from your set of friends? (If so why?)
17. Have you ever used Facebook to begin a long-distance relationship with someone you could not realistically meet? (If so, tell us about it.)
18. Have you ever considered ending a friendship/relationship over their Facebook behavior? (What did they do to make you consider this?)
USAGE HABITS
19. Do you use Facebook using the mobile app or your PC?
20. Do you post regularly on the dashboard?
21. Do you click on posts that link to other websites?
PROCEDURE
22. Can you explain to me how to set the privacy settings?
23. Can you explain to me how to change the password?
MONEY
24. Would you agree to pay a subscription to use Facebook? If yes, how much would you consider a reasonable amount to pay? (If not, why?)
25. If the application for PC available from your browser was free, but the mobile app was not. Would you pay for it?
26. Suppose that the free access to Facebook was limited in time, information you can access or which version of the app you can use. Which of these functionalities would have to be excluded from the free version for you to be interested in the subscription? Why that Specific one?
27. If Facebook would pay you in exchange for you performing tasks or taking surveys, would you be interested in them? (If yes, for how much? If the tasks could be considered unethical, would you still do it?)
28. Suppose Facebook will become a subscription service starting from tomorrow and you decide not to pay. What should Facebook do with your profile and data?
INFORMATION
29. When you read something that you find interesting, do you share it?(What motivates you to share it? Are you likely to share something without reading it?)
30. Is the information on Facebook more or less reliable than other sources? (For what reason?)
31. What is inappropriate information for Facebook? (Is there any information that should never reach Facebook? Should Facebook be used as a news source?)
PROCEDURE
32. Can you explain to me how create a post and tag someone into it?
33. Can you explain to me how to find friends that have no mutual friends?
ETHICS
34. FB censures some photos and posts if their content is signaled as inappropriate. Do you think this is correct? Where should the line be drawn between censure and freedom?
35. Recently FB has censored pictures of women breastfeeding even if the breast was not visible? Why do you think they do this? Should they be allowed to?
36. Recently FB workers admitted to routinely suppressing conservative content, do you feel they did anything wrong? (Why or why not?)
37. Should FB play a role in limiting/removing hate speech from the site? Is it ethical if they do?
38. Terrorist groups are known to have very active social media presences. Suppose Facebook began submitting information from all profiles to the government for help in tracking these groups. Would you be okay with this? Why?

TABLE 2
Machine learning features grouped by physiological and voice signal.

Signal	Features
EDA	- mean tonic - phasic AUC - phasic min, max, mean, sum peaks amplitudes
BVP	- min, max, sum peaks amplitudes - mean peak amplitude (diff. between baseline and task)
HR	- mean, sd. deviation (diff. between baseline and task) - mean Mel Spectrogram
Voice	- mean MFCC (mean of the first 20 MFCC features) - mean Chromagram (mean of the 12 Chroma features)

a band-pass filter algorithm at different intervals [37]. The EDA signal consists of a tonic component (i.e., the level of electrical conductivity of the skin) and a phasic one representing phasic changes in electrical conductivity or skin conductance response, or SCR [93]. We extract the two components using the `cvxEDA` algorithm [94].

After signals pre-processing, we extracted the features presented in Table 2, which we use to train our classifiers. We select biofeedback features based on previous studies using the same signals [26], [35], [74] and we choose audio features according to recommendations from the specialised literature [55], [91]. For the audio signal, we use the Python package `Librosa` [95]⁶.

We address the problem of predicting user engagement using machine learning classifiers. In line with previous research on biometrics [26], [35], [39], [74], we chose popular algorithms—i.e., Naive Bayes (NB), C4.5-like decision trees (DTree), Support Vector Machines (SVM), Multi-layer Perceptron for neural network (MLP), and Random Forest (RF).

3.6 Analysis Procedure

The analysis procedure aims at answering the three RQs. For each user, we use the biometrics gathered in the interview phase as input features for the different classifiers. We first consider solely biofeedback features (RQ1), then voice features (RQ2) and finally their combination (RQ3).

In line with previous research [26], [35], we target a *binary* classification task. In particular, we distinguish between *positive* vs *negative* valence, and *high* vs *low* arousal. As such, we exclude the *neutral* label from the gold standard and focus on more polarised values. Although this reduces our dataset, it also facilitates the separation between clearly distinguished emotional states⁷.

We evaluate our classifiers in the *hold-out* setting. Therefore, we split the gold standard into train (70%) and test (30%) sets using a stratified sampling strategy, which allows having a balanced set of instances from the different classes in both sets. For each algorithm, we search for the optimal hyper-parameters [96], [97] using leave-one-out cross validation on the train set—i.e., the recommended approach for small training sets [98] such as ours. The resulting model is then evaluated on the test set to assess its performance on unseen data and avoid overfitting. We repeat this process 10 times with different splits of the train and test sets to further

increase the validity of the results. The performance is then evaluated by computing the mean for precision, recall, F1-measure, and accuracy over the different runs. This setting is directly comparable to the one implemented by previous work [26], [35], which includes data from the same subject in both training and test sets.

The process outlined above is repeated with a maximum of eight different settings, based on the three following data pre-processing steps, aimed at improving the performance of the machine learning algorithms without losing validity of the results:

- **Standard Scaling:** the features in the training set are standardised so that their distribution will have a mean value 0 and standard deviation of 1. The standardisation parameters from the training set are then applied to scale the test set. This way, information from the test set (i.e., its standard deviation) is not passed to the training set, which could bias the learning process. Standard scaling is essential for machine learning algorithms that calculate distances between data, in our case SVM and MLP. If not scaled, the feature with a higher value range starts dominating when calculating distances. Scaling should not affect rule-based algorithms that consider each feature separately, and are not affected by monotonic transformations of the variables, such as standard scaling. Standard scaling is performed using the `StandardScaler` available in `scikit-learn`.
- **Balancing:** Synthetic Minority Oversampling Technique (SMOTE) is a traditional data augmentation technique applied to train machine learning models in case of class imbalance [99]. In case of class imbalance, machine learning algorithms tend to perform poorly on the minority class, as they do not have a sufficient amount of examples to learn from and build a fair classification model. To overcome this issue, SMOTE creates synthetic examples of the minority class, in our case based on the *k*-nearest neighbour algorithm. To prevent data leakage, SMOTE is applied solely to the training set, therefore the test set does not contain synthetic data. Balancing is performed using the SMOTE implementation in the `imblearn` Python package.
- **Imputation:** data imputation is normally adopted when some features have missing data. In our case, we miss voice feature data for 66 arousal vectors and 60 valence vectors. However, we can infer (*impute*) the data by using the corresponding biofeedback features. Imputation is performed using *k*-nearest neighbors approach implemented in the `KNNImputer` from `scikit-learn`.

4 EXECUTION AND RESULTS

The data were initially gathered from 31 participants. Interviews lasted 18 minutes on average. We discarded the data from those subjects for which data were largely incomplete, or that appeared to have a low degree of standard deviation (i.e., lower than one) in valence and arousal. Although these subjects may, in principle, have had little variations in their actual emotions, they can be considered outliers with

6. <https://librosa.org>

7. Preliminary experiments were performed considering a 3-label problem, but the number of vectors resulted too small to achieve acceptable results. This aspect is further discussed in Sect. 6.

respect of the rest of the subjects. As data are treated in aggregate form, and given the limited number of data points, including these outliers could have introduced undesired noise. We also discarded data whenever some inconsistency was observed through the different pre-processing steps. For example, our protocol required manual annotations of the timestamps during the interview, which sometimes led to not plausible timestamps. In these cases, we had to remove the data. At the end of this process, we produced the feature vectors and associated labels for valence and arousal (776 vectors in total from 21 subjects). The scatter plot for the two dimensions is reported in Fig. 1. The normalised range of the labels, evaluated by means of k-means clustering as explained in Sect. 3.6, is as follows. For valence we have: $[-4.94, -1.03]$ *negative*; $[-1.03, 2.52]$ *neutral*; $[2.52, 5.31]$ *positive*. For arousal we have: $[-4.8, 0.308]$ *low*; $[0.308, 3.57]$ *neutral*; $[3.57, 7]$ *high*.

TABLE 3
Label distribution in the gold standard for biofeedback feature vectors and for experiments using imputation.

Arousal			Valence		
High	Low	Neut.	Positive	Negative	Neut.
245 (66%)	191 (44%)	340	345 (79%)	89 (21%)	342

TABLE 4
Label distribution in the gold standard for voice feature vectors and combined feature vectors without imputation.

Arousal			Valence		
High	Low	Neut.	Positive	Negative	Neut.
159 (43%)	211 (57%)	303	300 (80%)	74 (20%)	270




Fig. 1. Scatter plots of normalised valence and arousal according to the self-assessment questionnaire.

As our goal is to discriminate between high (positive) and low (negative) arousal (valence), we removed all the

items for which the label resulted *neutral* for the dimensions, based on the participant’s answers. Therefore, our gold standard includes only the vectors labelled as high (positive) or low (negative) and we model our problem as a binary classification task. Table 3 reports the gold standard dataset with valence and arousal distribution, when considering biofeedback features (for RQ1). Voice feature vectors corresponding to each biofeedback vector could not be identified for part of the gold standard items, as the audio recording was not reliable for some subjects. Therefore, the gold standard dataset for audio only (RQ2) and for combined features (RQ3) without imputation is a subset of the original gold standard, and is reported in Table 4.




Fig. 2. Box plots of normalised valence and arousal for each group of questions, according to the self-assessment questionnaire.

4.1 Descriptive Statistics

In the following, we report descriptive statistics on the data. Table 5 reports the ranges of valence and arousal, according to the self-assessment questionnaire. We report both original values and normalised ones (“norm”, in the table). We see that, overall, users tend to give high scores both for arousal and valence (both averages are above 7), indicating that the interview is generally perceived as *positively engaging*. Although they used the whole 1 to 10 scale for both dimensions, indicating that the interview appeared to cover the whole range of emotions, we see that the standard deviation is not particularly large, especially for valence. Indeed, considering the 1-10 scale, the value of standard deviation (Std. Dev. in Table 5) indicates that around 68% of the subjects gave score $\in [6-9]$ for valence, and $\in [5-9]$ for arousal. This indicates that subjects tended to report scores around the average, and that apparently most of the interview triggered a similar level of engagement.

To gain more insight, it is useful to look at the reported engagement for each question⁸. Figure 2 reports the box plots for valence and arousal for each question, divided by question group. We see that questions related to *privacy*, *ethics* and *usage habits* tend to create more (positive) arousal

⁸. The statistics in this case consider solely those subjects that responded to *all* questions, i.e., 10 in total.

TABLE 5
Descriptive statistics of the reported engagement.

	Valence	Valence (norm)	Arousal	Arousal (norm)
Average	7.23	1.90	7.06	2.13
Minimum	1	-4.94	1	-4.8
Maximum	10	5.31	10	7
Std. Dev.	1.47	1.58	2.17	2.17

in average, while questions related to *procedures* are associated to more neutral values of arousal and valence (i.e., closer to 0 in the plot). Interestingly, questions related to *relationships* show the largest variation in terms of arousal and valence (the box-plot appears larger), indicating that this is a sensitive topic for the users, leading to more polarised scores in terms of emotional dimensions. The maximum average valence, instead, is observed for questions related to *ethics*.

4.2 RQ1: To what extent can we predict users' reported engagement using biofeedback measurements and supervised classifiers?

In Table 6 we report the performance of the different classifiers in terms of their precision, recall, F1-measure and accuracy, and considering the best configurations. Specifically, for each metric, we report the mean over the ten runs of the hold-out train-test procedure, i.e., the macro-average. This choice is in line with consolidated recommendations from literature on classification tasks using machine learning [100]. Specifically, using macro-averaging is recommended with unbalanced data as ours, as it emphasizes the ability of a classifier to behave well also on categories with fewer training instances on specific classes. For each classifier, we report its best performance and the associated configuration options in terms of data balancing and scaling.

We see that the best performance (in **bold**) for valence and arousal are achieved by Random Forest (RF), when applying balancing and standard scaling for valence (**Bal.** and **Scale** set to Y in the table), and balancing alone for arousal. The worst performing algorithm is Naive Bayes (NB), regardless of the configuration. A general pattern cannot be identified about the influence of the different configuration options on the performance of the algorithms. However, for RF, we observe that compensation for class imbalance has a positive impact for valence, which is also the dimension characterised by a limited number of negative data points (see Table 3), thus confirming the effectiveness of SMOTE when class imbalance is present.

In Table 7, we report the result of the two best algorithms with the best configurations, and we compare them with a baseline. Following previous research on sensor-based emotion recognition in software development [26], we select as baseline the trivial classifier always predicting the majority class, that is *high* for arousal and *positive* for valence. For the sake of completeness, we also report accuracy even if its usage is not recommended in presence of unbalanced data as ours.

For valence, the RF classifier distinguishes between negative and positive emotions with an F1 of 0.63, thus obtaining an increment of 40% with respect to the baseline. Furthermore, we observe an improvement in precision of 58% (from 0.40 of the baseline to 0.63 of RF) and 28% in

TABLE 6
Performance of the different algorithms with the different configurations in terms of data balancing with SMOTE (Bal.) and in terms of standard scaling (Scale) when using **biofeedback** features. Y = Yes, the configuration option is applied; N = No, the configuration option is not applied.

Algorithm	Options		Valence			
	Bal.	Scale	Prec	Rec	F1	Acc
SVM	Y	Y	0.601	0.617	0.603	0.717
MLP	N	N	0.614	0.581	0.582	0.773
DTree	N	Y	0.576	0.569	0.569	0.731
NB	N	Y	0.550	0.525	0.497	0.773
RF	Y	Y	0.633	0.635	0.633	0.758

Algorithm	Options		Arousal			
	Bal.	Scale	Prec	Rec	F1	Acc
SVM	N	Y	0.596	0.595	0.592	0.600
MLP	Y	Y	0.582	0.581	0.580	0.588
DTree	N	N	0.635	0.634	0.631	0.636
NB	N	N	0.514	0.510	0.468	0.544
RF	Y	N	0.658	0.656	0.654	0.660

TABLE 7
Performance of the best classifiers based on F1, using EDA, BVP, and HR features with respect to majority class baseline classifier. Improvement over the baseline is also shown.

	Precision	Recall	F1	Accuracy
Valence				
Random Forest	0.63	0.64	0.63	0.76
Baseline	0.40	0.50	0.45	0.79
Improvement	0.23 (58%)	0.14 (28%)	0.18 (40%)	-0.03 (-4%)
Arousal				
Random Forest	0.66	0.66	0.65	0.66
Baseline	0.28	0.50	0.36	0.56
Improvement	0.38 (135%)	0.16 (32%)	0.29 (81%)	0.1 (18%)

recall (from 0.50 to 0.64). These results indicate that the classifiers' behavior is substantially better than the baseline classifier that always predicts the positive class.

As for arousal, we observe a comparable performance. The RF classifier distinguishes between high and low activation with an F1 of 0.65, representing an improvement of 81% over the baseline (0.36). Again, the classifier substantially outperforms the baseline with an improvement of 135% for precision (from 0.28 to 0.66) and 32% for recall (from 0.50 to 0.66). The improvement with respect to the baseline is particularly high for arousal in terms of precision since arousal data (cf. Table 3) are more balanced. Therefore, the baseline that always predicts the positive class is inherently less effective, with a precision of 0.28.

4.3 RQ2: To what extent can we predict users' engagement using voice analysis and supervised classifiers?

Table 8 reports the comparison of the performance between the different machine learning algorithms, considering the best configurations. Differently from the biofeedback case, here we apply also data imputation (**Imp.** column) as configuration option, by synthesising data for vectors with missing voice data, based on the corresponding biofeedback vectors. Therefore, the gold standard considered when **Imp.** is set to Y (Yes) is analogous to the one used for biofeedback and reported in Table 3. When **Imp.** is set to N (No), the gold standard is the one in Table 4.

This is an important aspect to notice for at least two reasons: (1) the two gold standards have slightly different distributions, and thus the default baselines for compari-

TABLE 8

Performance of the different algorithms with the different configurations in terms of data balancing with SMOTE (Bal.), standard scaling (Scale) and Imputation (Imp.) when using **voice** features. Y = Yes, the configuration option is applied; N = No, the configuration option is not applied.

Algorithm	Options			Valence			
	Bal.	Scale	Imp.	Prec	Rec	F1	Acc
SVM	N	Y	N	0.703	0.680	0.687	0.810
MLP	N	Y	N	0.688	0.680	0.680	0.802
DTree	Y	N	N	0.636	0.662	0.644	0.756
NB	N	Y	N	0.611	0.652	0.612	0.698
RF	Y	N	N	0.709	0.693	0.697	0.817

Algorithm	Options			Arousal			
	Bal.	Scale	Imp.	Prec	Rec	F1	Acc
SVM	Y	Y	N	0.710	0.710	0.708	0.714
MLP	Y	Y	N	0.703	0.703	0.701	0.706
DTree	Y	N	Y	0.634	0.635	0.634	0.639
NB	Y	N	Y	0.642	0.643	0.640	0.643
RF	Y	Y	Y	0.684	0.683	0.681	0.685

son will differ. In particular, when using imputation the baselines are the same as the one used in Table 7. When imputation is not used, the baselines need to be recomputed, and in particular the majority class baseline for arousal will always predict the low class, as this is the most frequent in Table 4; (2) the usage of imputation in a real-world context assumes that, although solely voice data are used for classification, biofeedback data are collected anyway, so the practical advantage, both economic and logistic, is limited.

The best performance (in **bold**) is achieved by Random Forest (RF) for valence, and by Support Vector Machines (SVM) for arousal, while NB again remains far behind the other algorithms, regardless of the applied configuration options.

For valence, the best performance is achieved by RF when balancing is applied and scaling and imputation are *not* applied. The performance is higher than the ones obtained with biofeedback features (F1 = 0.70 vs F1=0.63 for biofeedback). This indicates that voice features alone appear to be particularly effective in discriminating the quality of the engagement (positive or negative valence), thereby confirming that our set-up is adequate for capturing the so called *emotional prosody* [25], which reveals the sentiment of the speaker.

Lower performance is achieved with imputation for most of the algorithms (Table 8 reports the best performing configurations, and Imp. = N for most of the high-performing cases). This suggests that biofeedback and voice capture different independent aspects of valence, as using biofeedback information to enrich voice data is not effective.

For arousal, the best performance is achieved by SVM when applying balancing and scaling, again without imputation. The performance is higher than the best one obtained with biofeedback features (F1=0.71 for SVM vs F1=0.65 for RF, cf. Table 6).

Overall, our findings suggest that voice features represent a valid, and even more effective alternative to biofeedback for the emotion recognition during requirements elicitation. In fact, our classifiers' performance demonstrate that in absence of biofeedback information, both valence and arousal can be successfully predicted with voice-only features.

TABLE 9

Performance of the best classifiers, according to F1, using voice features with respect to majority class baseline classifier. Improvement over the baseline is also shown.

	Precision	Recall	F1	Accuracy
Valence				
Random Forest	0.71	0.69	0.70	0.82
Baseline	0.40	0.50	0.45	0.80
Improvement	0.31 (77%)	0.19 (39%)	0.25 (55%)	0.02 (2%)
Arousal				
SVM	0.71	0.71	0.71	0.71
Baseline	0.29	0.50	0.37	0.57
Improvement	0.42 (145%)	0.21 (42%)	0.34 (91%)	0.14 (25%)

TABLE 10

Performance of the classifiers with the different configurations, considering biofeedback and voice features **combined**.

Algorithm	Options			Valence			
	Bal.	Scale	Imp.	Prec	Rec	F1	Acc
SVM	Y	Y	N	0.723	0.718	0.720	0.826
MLP	Y	Y	N	0.692	0.705	0.696	0.803
DTree	Y	Y	N	0.621	0.650	0.626	0.734
NB	Y	Y	N	0.617	0.663	0.617	0.701
RF	Y	Y	N	0.695	0.693	0.692	0.808

Algorithm	Options			Arousal			
	Bal.	Scale	Imp.	Prec	Rec	F1	Acc
SVM	N	Y	N	0.711	0.701	0.700	0.712
MLP	Y	Y	N	0.690	0.688	0.687	0.694
DTree	Y	Y	N	0.665	0.664	0.660	0.665
NB	Y	N	N	0.631	0.631	0.627	0.632
RF	Y	Y	N	0.726	0.724	0.723	0.728

To have additional insights, Table 9 compares the result of the best algorithms, with respect to the majority class baselines. For valence, we seen an increase of 77% in terms of precision, 39% for recall and 55% for F1. For arousal, the increase in performance is again higher, with 145% for precision, 42% for recall and 91% for F1.

4.4 RQ3: To what extent can we predict users' engagement by combining voice and biofeedback features?

Table 10 reports the comparison of the performance of the different algorithms with their best configurations when using voice and biofeedback features combined.

General trends are analogous to those observed when features are treated in separation, with the best performance achieved by RF and SVM, and lower performance by DTree and NB.

TABLE 11

Comparison of the performance for all the algorithms, considering their best configurations for the different feature combinations. In **bold**, we report the best performance for each algorithm considering F1.

Algorithm	Features	Valence			Features	Arousal		
		Prec	Rec	F1		Prec	Rec	F1
SVM	biofeedback	0.601	0.617	0.603	biofeedback	0.596	0.595	0.592
	voice	0.703	0.680	0.687	voice	0.710	0.710	0.708
	combined	0.723	0.718	0.720	combined	0.711	0.701	0.700
MLP	biofeedback	0.580	0.583	0.578	biofeedback	0.582	0.581	0.580
	voice	0.688	0.680	0.680	voice	0.703	0.703	0.701
	combined	0.692	0.705	0.696	combined	0.690	0.688	0.687
DTree	biofeedback	0.576	0.569	0.569	biofeedback	0.635	0.634	0.631
	voice	0.636	0.662	0.644	voice	0.634	0.635	0.634
	combined	0.621	0.650	0.626	combined	0.665	0.664	0.660
NB	biofeedback	0.550	0.525	0.497	biofeedback	0.514	0.510	0.468
	voice	0.611	0.652	0.612	voice	0.642	0.643	0.640
	combined	0.617	0.663	0.617	combined	0.631	0.631	0.627
RF	biofeedback	0.633	0.635	0.633	biofeedback	0.658	0.656	0.654
	voice	0.709	0.693	0.697	voice	0.684	0.683	0.681
	combined	0.695	0.693	0.692	combined	0.726	0.724	0.723

TABLE 12
Performance of the best classifiers, according to F1, using combined features with respect to majority class baseline classifiers. Improvement over the baseline is also shown.

	Precision	Recall	F1	Accuracy
<i>Valence</i>				
SVM	0.72	0.73	0.72	0.83
Baseline	0.40	0.50	0.45	0.79
Improvement	0.32 (80%)	0.23 (46%)	0.27 (60%)	0.04 (5%)
<i>Arousal</i>				
Random Forest	0.73	0.72	0.72	0.73
Baseline	0.29	0.50	0.37	0.57
Improvement	0.44 (150%)	0.22 (45%)	0.35 (95%)	0.16 (28%)

The best performance for valence (in **bold**) is achieved by SVM, when applying SMOTE and scaling, and without imputation (F1=0.72). Instead, the best performance for arousal is obtained by RF with the same configuration (F1=0.72). These results suggest that voice and biofeedback features can play complementary roles in engagement recognition, and can lead to the best classification results compared to those obtained when using only one source of data, when specific algorithms are selected. These observations are confirmed by Table 11, where the best performance of each algorithm are compared, considering different feature combinations. The table shows that the best performance in terms of F1 (in **bold**) is achieved either using voice feature alone, or by combining the features.

As for the other cases, in Table 12 we compare the performance for the best algorithms with the majority class baselines. In all cases, the improvement for precision, recall and F1 is always greater or equal to 45%, thus confirming that the combination of voice and biofeedback features allows obtaining fine-grained distinction of classes for both valence and arousal, which cannot be achieved without combining the features.

5 DISCUSSION

The main take-away messages of this study are:

- users’ interviews are activities that can trigger positive engagement in the involved users;
- different levels of engagement are experienced depending on the topic of the question, with topics such as *privacy*, *ethics* and *usage habits* leading to higher engagement, and *relationships* leading to larger variations of engagement;
- by combining biofeedback features into vectors and by training the Random Forest (RF) algorithm, it is possible to predict the engagement in a way that outperforms a majority-class baseline, with F1-measure of 63% for valence and 65% for arousal;
- using voice features only when training Support Vector Machines (SVM) and RF, performance are increased. Engagement can be predicted through voice features alone with F1-measure 70% (valence, RF) and 71% (arousal, SVM)
- the combination of biofeedback and voice features maximises the performance, with F1 measure 72% (valence) and 72% (arousal). In this case, the best performance is achieved with RF for arousal, and with SVM for valence.

In the following sections, we discuss our results in relation to existing literature and outline possible applications and timely avenues of research that are enabled by the current study.

5.1 Engagement and Topics

Our descriptive statistics indicate that users experienced different levels of engagement with respect to the topic of a question. Specifically, our participants reported a positive attitude when discussing privacy, ethics, and usage habits. Concerning privacy and ethics, these topics were selected on purpose to trigger higher engagement. Given the raising interest in these two fields, especially in relation to Facebook and online communities in general—e.g., [101]—the obtained results are not surprising. Concerning *usage habits*, we expected to see lower values of arousal. Questions regarding usage habits were asked at the beginning. Therefore, the observed high arousal could result from the excitement of the new experience. However, we observed that question 19 (also about usage habits, but asked later) had the highest average arousal (3.6 in normalised values, while the average for usage habits questions is 2.8) and valence (3.2 vs. 2.5)⁹. We therefore argue that talking about usage habits triggers positive engagement. This indicates that users generally like the platform and are interested in sharing their relation with it. Qualitative analysis of the audio of the actual answers, not performed in this study, can further clarify these aspects. Overall, these results show that 1) users’ interviews elicit emotions and engagement, with varying degrees of reactions depending on the topic; and 2) some topics are perceived as more engaging than others.

5.2 Performance Comparison with Related Studies

According to the theoretical model of affect described in Sect. 2, in this study we use emotions as a proxy for engagement. Specifically, we operationalize emotions along the valence and arousal dimensions of the Circumplex Model of affect [23], which we recognize using biometrics and voice. Using machine learning, we are able to classify engagement of users during requirements interviews by distinguishing between positive and negative valence and high and low arousal. We experimented with different experimental setting—i.e., data balancing using SMOTE, data scaling, and data imputation (for voice data only).

As for our results using biofeedback features, a direct comparison is possible with the performance achieved in the empirical study by Girardi et al. [26], as we use the same device (i.e., Empatica E4 wristband) and include the same metrics for EDA, BVP, and HR. Our classifier performance for arousal (F1 = 0.63) and valence (F1 = 0.65) outperforms the one they obtain using Empatica—i.e., 0.55 for arousal and 0.59 for valence. They report a slightly better performance, though still lower than ours, when including also the signals gathered using an EEG helmet (F1 = 0.59 for arousal and F1 = 0.60 for valence). Müller and Fritz [35] report an accuracy of 0.71 for valence, using a combination of features based on EEG, HR, and pupil size captured by an eye-tracker. Their goals are different from ours, as they

9. results for each individual question not shown in the paper

deal with programming tasks, and neither voice nor active stimulation of emotions were present in these related works. Our better performance can be specifically linked to the task of interviewing in which voice is not only used as a feature for predicting emotions, but also as mean for expressing them [102], [103]. The simple act of vocalizing can be regarded as an explicit, although not necessarily voluntary, expression of emotion. This could affect other biometric aspects, improving the performance of our classifiers also when biofeedback-based features are the only predictors.

Previous studies in affective computing report much higher performance—e.g., accuracy of 0.97 for arousal [104], [105], [106] and 0.91 for valence [47]—for tasks in which emotions were recognized from standard stimuli (e.g., videos in the DEAP dataset [39]). These studies rely on high-definition EEG helmets [104], [105], [106] and facial electrodes for EMG [47] which are invasive and cannot be used during real interviews with users or in remote interviews.

Considering our results when using voice features, our approach achieves higher performance with respect to biofeedback for both arousal ($F1 = 0.71$) and valence ($F1 = 0.70$). The results of the voice-based classifier we trained and tested in the scope of this study are in line with several machine learning-based state-of-the-art approaches in the broader field of speech emotion recognition [52]. The model relying on voice-based features paves the way to future replications for *in vivo* studies that do not require the use of wearable sensors.

As far as the combination of biofeedback and voice is concerned, our classifier shows comparable performance with respect to a deep-learning based approach recently proposed by Aledhari et al. [107]. As in our study, they use the Empatica E4 wristband for collecting biofeedback and report an accuracy of 85% on test set and 79% in the validation set for the recognition of emotional valence (our best accuracy is 83%).

5.3 Implications for Research and Practice

We consider this an exploratory study which helps us and the community to have a first understanding of engagement in user interviews and the potential usage of biofeedback devices and voice analysis in this context. However, we argue that our results, once consolidated and confirmed by further studies, can have multiple applications and can open new avenues of research.

5.3.1 Applications in User Feedback

In user interviews similar to those staged in our experiment, biofeedback and voice information can be used to better investigate possible discrepancies between user engagement and the reported relevance of features. The ability to do so can facilitate requirements prioritization tasks the same way sentiment analysis does when applied to textual user feedback [80]. Furthermore, the usage of these technologies can be extended to identify the engagement of the user *on-the-fly*—i.e., during the interview—and help analysts steering the flow of the conversation. These applications, which support human analysts in their activity, become particularly important when *artificial agents* are used to

elicit feedback or provide customer support, as shown by related research on voice analysis for call centers [68], [69]. In these context, the detection of negative emotions is used to understand when a human operator needs to replace an artificial one, because the latter is irritating the customer. Our work opens to further applications on emotion-aware, voice-based chatbots for user interviews.

5.3.2 The Role of Voice

The introduction of voice features is particularly crucial. Biofeedback needs to be locally acquired with specialised devices such as Empatica E4, which (i) costs about \$1,690.00 at the time of writing; (ii) needs to locally register the different signals; (iii) does not remotely send the signal in an automated manner; (iv) can raise privacy concerns. Therefore, their usage is realistic only during face-to-face interviews, in which a certain level of mutual trust can be achieved and all data can be acquired locally. Instead, the analysis of voice is particularly appropriate in remote communication scenarios—involving either human or artificial agents—which are increasingly common due to the COVID-19 pandemic. Voice is voluntarily produced and transmitted by users, and can be remotely recorded and processed without resorting to specialised devices, with evident cost savings. The cost reduction extends the applicability of the idea to large-scale scenarios. With voice analysis, automated user feedback campaigns become feasible, and companies can improve automated A/B testing of web apps or pages. Specifically, they can ask multiple users to interact with different versions of an interface, and speak-up their reflections on the experience. The recording and the analysis of the engagement can be used to facilitate the identification of preferred versions, appreciated features, or relevant interaction problems.

5.3.3 Applications in RE and Software Engineering.

In the case of more classical requirements elicitation interviews [9], [11], the usage of biometrics can support these activities to improve the analyst's ability to create a trustworthy relationship with the customer, and improve the quality of the interview and the collected data. In this context, it is relevant to extend the work to identify the customer's frustration, which often corresponds to the first step to create mistrust in the analyst [108]. Frustration can be detected using biofeedback by analyzing the changes in the heart-rate, temperature, and other vitals [109], [110], [111], [112] and used to warn the analyst. Furthermore, frustration is strictly related with stress, which can be detected in voice signals through Teager energy operator (TEO)-based features [113], [114].

Overall, we argue that the analysis of voice, with its relative cost-effectiveness, can be broadly applied not only to RE, but to all software engineering scenarios in which conversations are central (e.g., SCRUM stand-up meetings, synchronous code reviews, pair-programming [115]) to investigate the emotional side of these human-intensive activities that have a relevant impact on the development, but are currently ephemeral.

5.3.4 Tacit Knowledge

It is worth noting that the improved performance obtained with voice features, and the lower cost of the approach, do not rule-out biofeedback. We have shown that the best performance are actually obtained with a combination of both types of features. In addition, biofeedback captures involuntary body signals that the speaker cannot fully control, while voice tone can, to a certain extent, be manipulated to deceive [116]. Biofeedback can reveal a more faithful representation of emotions, and one can compare discrepancies between emotion prediction with biofeedback and with voice to identify situations in which what the voice “tells” is different from what the speaker “feels.” This can happen in requirements elicitation interviews, which can involve controversial political aspects [117], or domain experts who need to be interviewed to gather process-related information but may be reluctant to share their knowledge [118]. Therefore, the results of this research can further be applied to tackle the issue of tacit knowledge in requirements engineering [118], [119], [120].

6 THREATS TO VALIDITY

In this section, we report the main limitations of our study and how we address them. The order in which we report the threats follows Wohlin et al. [121] suggestion regarding exploratory research.

Internal validity. Threats to internal validity deal with confounding factors that can influence the results of a study. We collected data in a laboratory setting. Factors existing in our settings, such as the presence of the experimenter, can influence the emotional status as the participants [122]. Establishing a trust-based rapport with the participants in a relaxed setting is crucial to mitigate these threats. Thus, we invited the participant to wear the wristband when entering in the room, before the actual interview started, in order to get acquainted with the device, settings, and the presence of the experimenter. Furthermore, self-assessment questionnaires were filled *immediately after* the interview. This choice was driven by the need to preserve a realistic interview context. However, with this design, the engagement is *recalled* by the subject and not reported in the moment in which it emerged. Therefore, discrepancies can occur between the feeling of engagement and its rationally-processed memory. Similarly, to maintain a realistic settings, we did not perform pre-interviews to assess the participants’ mood (i.e., the presence of a long-lasting emotion) nor their personality traits. We acknowledge that an emotionally-charged event in the life of a participant, occurring before the interview, can impact the results.

In our design, we consider a *hold-out* setting for validation in line with previous studies [28], [76], [123]—i.e., a subject could contribute with different vectors to training and test sets; however, in practice, this can happen only if the same subject is interviewed multiple times. This choice is driven by the limited number of subjects which does not allow to have effective predictions in a more realistic *leave-one-subject-out* (LOSO) setting. For example, [28] perform an *in vivo* study involving 21 developers using the LOSO evaluation settings, showing an $F1=0.46$ for valence.

Construct validity. This threat refers to the reliability of the operationalization of the study constructs. Our study may suffer by threats to construct validity in capturing emotions using self-reports. To address this issue, we performed data quality assurance and excluded participants who did not show engagement with the task (e.g., who provided always the same score or scores with overall low standard deviation). We believe that the designed interview script is sufficiently representative of typical users’ interviews in terms of triggered engagement. However, we did not pause the interview when addressing a new topic without giving the respondents the chance to relax and reduce their emotional agitation (e.g., by watching a relaxing video). Therefore, the order in which we presented the topics can have an impact on the result—i.e., for some topics presented later in the interview we may be observing an *emotional spillover* due to emotional charge introduced when discussing earlier topics. Concerning the self-assessment questionnaire, this was adapted from previous studies in software engineering [76], and the users gained confidence with self-assessment in the initial emotion triggering activity.

We exclude neutral cases from our dataset. This choice is driven by the need to simplify the prediction problem. Considering the exploratory nature of the study, we decided to focus first on more polarised emotional states, which are assumed to be clearly distinguishable. Preliminary experiments including neutral cases showed an $F1=0.45$ for both valence and arousal using RF with biofeedback features, in a hold-out setting like the one adopted in the rest of the experiments of our paper.

Conclusion validity. The validity of our conclusions relies on the robustness of the machine learning models. To mitigate any threat arising from having a small dataset, we ran several algorithms addressing the same classification task. In all runs, we performed hyperparameters tuning as recommended by state-of-the-art research [124]. Following consolidated guidelines for machine learning, we split our data into train-test subsets. The training is performed using cross-validation and the final model performance is assessed on a hold-out test set. The entire process is repeated ten times for each algorithm, to account for random variations in the data. Moreover, our classifiers configuration included scaling and data balancing techniques.

External validity. The generalizability of our results is limited by the amount of subjects (and associated data points) who took part in the study. Although with some imbalance, our sample includes multiple ethnic groups and genders to account for physiological differentiation [87]. Further replications with a confirmatory design should engage more participants, and consider balance between ethnicity, culture, age, and gender to account for the differences in emotional reactions due to these aspects.

As for the topic of the interviews, we selected features from a commonly-used social media app for which no particular expertise is needed. We explicitly base our questionnaire on the Facebook social media platform due to its familiarity for both interviewers and interviewees. However, the questions are generic enough to cover functionalities of other platform in the same space. The interviews were performed in 2016 before the Facebook–Cambridge Analytica data scandal was reported by *The New York Times*

in December 2016 [125]. Up to that point, the only report of profiles harvested in Facebook was published by the Guardian in December 2015 [126], and Facebook was not considered a controversial platform.

7 CONCLUSION AND FUTURE WORK

This paper presents the first study about engagement prediction in user interviews. In particular, we show that it is possible to predict the positive or negative engagement of a user during an interview about a product. This can be achieved through the usage of *biofeedback* measurements acquired through a wristband, the analysis of voice through audio processing, and the application of supervised machine learning. In particular, we show that voice analysis alone can lead to sufficiently good results.

The study is exploratory in nature, and application of our results requires further investigation, and the acceptance of the non-intrusive, yet potentially undesired, biofeedback device. Nevertheless, we believe that the current work, with its promising results, establishes the basis for further research on engagement, and emotions in general, during the many human-intensive activities of system development. Among the future works, we plan to: (a) replicate the experiment with a larger and more representative sample of participants, utilizing a realistic validation setting—e.g., including neutral data; (b) complement our analysis with the usage of other emotion-revealing signals, such as facial expressions captured through cameras [40] and electroencephalographic (EEG) activity data [26], [35]; (c) tailor the study protocol for requirements elicitation interviews for novel products yet to be developed; (d) analyse emotions of interviewers and interviewee during the dialogue, to investigate if it is possible to identify the process of construction of trust; (e) investigate and compare the emotional footprint of different software development-related tasks. This can be done for example by looking at the difference between physiological signals of the multiple actors of the development process across different phases, such as of development, elicitation, testing, *etc*; (f) apply voice analysis to SCRUM meetings, focus groups, and other software engineering activities in which speech plays a primary role.

REFERENCES

- [1] M. Bano and D. Zowghi, "A systematic review on the relationship between user involvement and system success," *Information and Software Technology*, vol. 58, pp. 148–169, 2015.
- [2] Z. Bakalova and M. Daneva, "A comparative case study on clients participation in a traditional and in an agile software company," in *Proc. of the 12th Int. Conf. on product focused software development and process improvement*, 2011, pp. 74–80.
- [3] G. K. Hanssen and T. E. Fægri, "Agile customer engagement: a longitudinal qualitative case study," in *Proc. of the 2006 ACM/IEEE Int. symposium on Empirical software engineering*, 2006, pp. 164–173.
- [4] J. Heiskari and L. Lehtola, "Investigating the state of user involvement in practice," in *2009 16th Asia-Pacific Software Engineering Conf.* IEEE, 2009, pp. 433–440.
- [5] P. K. Murukannaiah, N. Ajmeri, and M. P. Singh, "Acquiring creative requirements from the crowd: Understanding the influences of personality and creative potential in crowd re," in *2016 IEEE 24th Int. Requirements Engineering Conf. (RE)*. IEEE, 2016, pp. 176–185.
- [6] E. C. Groen, N. Seyff, R. Ali, F. Dalpiaz, J. Doerr, E. Guzman, M. Hosseini, J. Marco, M. Oriol, A. Perini *et al.*, "The crowd in requirements engineering: The landscape and challenges," *IEEE software*, vol. 34, no. 2, pp. 44–52, 2017.
- [7] W. Maalej, M. Nayebi, T. Johann, and G. Ruhe, "Toward data-driven requirements engineering," *IEEE Software*, vol. 33, no. 1, pp. 48–54, 2015.
- [8] G. Williams and A. Mahmoud, "Mining twitter feeds for software user requirements," in *2017 IEEE 25th Int. Requirements Engineering Conf. (RE)*. IEEE, 2017, pp. 1–10.
- [9] D. Zowghi and C. Coulin, "Requirements elicitation: A survey of techniques, approaches, and tools," in *Engineering and managing software requirements*. Springer, 2005, pp. 19–46.
- [10] D. M. Fernández, S. Wagner, M. Kalinowski, M. Felderer, P. Mafra, A. Vetrò, T. Conte, M.-T. Christiansson, D. Greer, C. Lassenius *et al.*, "Naming the pain in requirements engineering," *Empirical software engineering*, vol. 22, no. 5, pp. 2298–2338, 2017.
- [11] A. Davis, O. Dieste, A. Hickey, N. Juristo, and A. M. Moreno, "Effectiveness of requirements elicitation techniques: Empirical results derived from a systematic review," in *14th IEEE Int. Requirements Engineering Conf. (RE'06)*. IEEE, 2006, pp. 179–188.
- [12] I. Hadar, P. Soffer, and K. Kenzi, "The role of domain knowledge in requirements elicitation via interviews: an exploratory study," *Requirements Engineering*, vol. 19, no. 2, pp. 143–159, 2014.
- [13] C. Palomares, X. Franch, C. Quer, P. Chatzipetrou, L. López, and T. Gorschek, "The state-of-practice in requirements elicitation: an extended interview study at 12 companies," *Requirements Engineering*, vol. 26, no. 2, pp. 273–299, 2021.
- [14] A. M. Aranda, O. Dieste, and N. Juristo, "Effect of domain knowledge on elicitation effectiveness: an internally replicated controlled experiment," *IEEE Transactions on Software Engineering*, vol. 42, no. 5, pp. 427–451, 2015.
- [15] A. Ferrari, P. Spoletini, and S. Gnesi, "Ambiguity and tacit knowledge in requirements elicitation interviews," *Requirements Engineering*, vol. 21, no. 3, pp. 333–355, 2016.
- [16] M. Bano, D. Zowghi, A. Ferrari, P. Spoletini, and B. Donati, "Teaching requirements elicitation interviews: an empirical study of learning from mistakes," *Requirements Engineering*, vol. 24, no. 3, pp. 259–289, 2019.
- [17] E. Guzman and W. Maalej, "How do users like this feature? a fine grained sentiment analysis of app reviews," in *2014 IEEE 22nd Int. requirements engineering Conf. (RE)*. IEEE, 2014, pp. 153–162.
- [18] Z. Kurtanović and W. Maalej, "On user rationale in software engineering," *Requir. Eng.*, vol. 23, no. 3, p. 357–379, Sep. 2018. [Online]. Available: <https://doi.org/10.1007/s00766-018-0293-2>
- [19] S. A. Scherr, C. Kammler, and F. Elberzhager, "Detecting user emotions with the true-depth camera to support mobile app quality assurance," in *2019 45th Euromicro Conf. on Software Engineering and Advanced Applications (SEAA)*. IEEE, 2019, pp. 169–173.
- [20] P. Mennig, S. A. Scherr, and F. Elberzhager, "Supporting rapid product changes through emotional tracking," in *2019 IEEE/ACM 4th Int. Workshop on Emotion Awareness in Software Engineering (SEmotion)*. IEEE, 2019, pp. 8–12.
- [21] P. Spoletini, C. Brock, R. Shahwar, and A. Ferrari, "Empowering requirements elicitation interviews with vocal and biofeedback analysis," in *2016 IEEE 24th Int. Requirements Engineering Conf. (RE)*. IEEE, 2016, pp. 371–376.
- [22] K. Taveter, L. Sterling, S. Pedell, R. Burrows, and E. M. Taveter, "A method for eliciting and representing emotional requirements: Two case studies in e-healthcare," in *2019 IEEE 27th Int. Requirements Engineering Conf. Workshops (REW)*. IEEE, 2019, pp. 100–105.
- [23] J. Russell, "Culture and the categorization of emotions," *Psychological Bulletin*, vol. 110 (3), pp. 426–450, 1991.
- [24] D. Girardi, A. Ferrari, N. Novielli, P. Spoletini, D. Fucci, and T. Huichapa, "The way it makes you feel predicting users' engagement during interviews with biofeedback and supervised learning," in *28th IEEE International Requirements Engineering Conference, RE 2020, Zurich, Switzerland, August 31 - September 4, 2020*, T. D. Breaux, A. Zisman, S. Fricker, and M. Glinz, Eds. IEEE, 2020, pp. 32–43. [Online]. Available: <https://doi.org/10.1109/RE48521.2020.00016>
- [25] T. W. Buchanan, K. Lutz, S. Mirzazade, K. Specht, N. J. Shah, K. Zilles, and L. Jäncke, "Recognition of emotional prosody and

- verbal components of spoken language: an fmri study," *Cognitive Brain Research*, vol. 9, no. 3, pp. 227–238, 2000.
- [26] D. Girardi, N. Novielli, D. Fucci, and F. Lanubile, "Recognizing Developers' Emotions while Programming," in *42nd Int. Conf. on Software Engineering (ICSE '20)*, May 23–29, 2020, Seoul, Republic of Korea, 2020. [Online]. Available: <https://doi.org/10.1145/3377811.3380374>
- [27] T. Fritz, A. Begel, S. C. Müller, S. Yigit-Elliott, and M. Züger, "Using psycho-physiological measures to assess task difficulty in software development," in *36th Int. Conf. on Software Engineering, ICSE '14, Hyderabad, India - May 31 - 7 June, 2014*, 2014, pp. 402–413. [Online]. Available: <https://doi.org/10.1145/2568225.2568266>
- [28] D. Girardi, F. Lanubile, N. Novielli, and A. Serebrenik, "Emotions and perceived productivity of software developers at the workplace," *IEEE Transactions on Software Engineering*, 2021.
- [29] A. Ferrari, T. Huichapa, P. Spoletini, N. Novielli, D. Fucci, and D. Girardi, "Replication Package: Using Voice and Biofeedback to Predict User Engagement during Requirements Interviews," 2021.
- [30] R. Cowie, N. Sussman, and A. Ben-Ze'ev, "Emotion: Concepts and definitions," in *Emotion-oriented systems*. Springer, 2011, pp. 9–30.
- [31] A. Ortony, G. Clore, and A. Collins, *The Cognitive Structure of Emotion*, 01 1988, vol. 18.
- [32] R. S. Lazarus, *Emotion and Adaptation*. Oxford University Press, 1991.
- [33] C. Yu, P. M. Aoki, and A. Woodruff, "Detecting user engagement in everyday conversations," in *INTERSPEECH 2004 - ICSLP, 8th International Conference on Spoken Language Processing, Jeju Island, Korea, October 4-8, 2004*. ISCA, 2004. [Online]. Available: http://www.isca-speech.org/archive/interspeech_2004/i04_1329.html
- [34] A. Barhenke, A. L. Müller, E. Brown, R. Seifer, and S. Dickstein, "Observed emotional and behavioral indicators of motivation predict school readiness in head start graduates," *Early Childhood Research Quarterly*, vol. 26, no. 4, pp. 430–441, 2011.
- [35] S. C. Müller and T. Fritz, "Stuck and frustrated or in flow and happy: Sensing developers' emotions and progress," in *37th IEEE/ACM Int. Conf. on Software Engineering, ICSE 2015, Florence, Italy, May 16-24, 2015*, pp. 688–699.
- [36] D. Grazier, X. Wang, and P. Abrahamsson, "Do feelings matter? on the correlation of affects and the self-assessed productivity in software engineering," *Journal of Software: Evolution and Process*, vol. 27, no. 7, pp. 467–487, 2015. [Online]. Available: <https://doi.org/10.1002/smr.1673>
- [37] F. Canento, A. Fred, H. Silva, H. Gamboa, and A. Lourenço, "Multimodal biosignal sensor data handling for emotion recognition," in *SENSORS*. IEEE, 2011, pp. 647–650.
- [38] J. Kim and E. André, "Emotion recognition based on physiological changes in music listening," *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 30, no. 12, pp. 2067–2083, 2008. [Online]. Available: <https://doi.org/10.1109/TPAMI.2008.26>
- [39] S. Koelstra, C. Muehl, M. Soleymani, J. Lee, A. Yazdani, T. Ebrahimi, T. Pun, A. Nijholt, and I. Patras, "DEAP: A database for emotion analysis using physiological signals," *IEEE Transaction on Affective Computing*, vol. 3, no. 1, pp. 18–31, 2012. [Online]. Available: <https://doi.org/10.1109/T-AFFC.2011.15>
- [40] M. Soleymani, S. Asghari-Esfeden, Y. Fu, and M. Pantic, "Analysis of EEG signals and facial expressions for continuous emotion detection," *IEEE Transaction on Affective Computing*, vol. 7, no. 1, pp. 17–28, 2016. [Online]. Available: <https://doi.org/10.1109/T-AFFC.2015.2436926>
- [41] D. Girardi, F. Lanubile, and N. Novielli, "Emotion detection using noninvasive low cost sensors," in *Seventh Int. Conf. on Affective Computing and Intelligent Interaction, ACII 2017, San Antonio, TX, USA, October 23-26, 2017*, 2017, pp. 125–130.
- [42] A. E. Kramer, *Physiological Metrics of Mental Workload: A Review of Recent Progress*, D. T. I. Center, Ed., 06 1990.
- [43] B. Reuderink, C. Muehl, and M. Poel, "Valence, arousal and dominance in the eeg during game play," *Int. Journal of Autonomous and Adaptive Communications Systems*, vol. 6, no. 1, pp. 45–62, 2013. [Online]. Available: <http://dx.doi.org/10.1504/IJAACS.2013.050691>
- [44] M. Li and B.-L. Lu, "Emotion classification based on gamma-band eeg," in *2009 Annual Int. Conf. of the IEEE Engineering in Medicine and Biology Society, Sep. 2009*, pp. 1223–1226.
- [45] W. Burleson and R. W. Picard, "Affective agents: Sustaining motivation to learn through failure and state of "stuck"," in *Social and Emotional Intelligence in Learning Environments Workshop.*, 8 2004.
- [46] A. Kapoor, W. Burleson, and R. W. Picard, "Automatic prediction of frustration," *Int. Journal Human-Computer Studies*, vol. 65, no. 8, pp. 724–736, Aug. 2007.
- [47] P. A. Nogueira, R. A. Rodrigues, E. C. Oliveira, and L. E. Nacke, "A hybrid approach at emotional state detection: Merging theoretical models of emotion with data-driven statistical classifiers," in *2013 IEEE/WIC/ACM Int. Conf. on Intelligent Agent Technology, IAT 2013*. IEEE, 2013, pp. 253–260.
- [48] J. Scheirer, R. Fernandez, J. Klein, and R. W. Picard, "Frustrating the user on purpose: a step toward building an affective computer," *Interacting with Computers*, vol. 14, pp. 93–118, 2002. [Online]. Available: <https://ieeexplore.ieee.org/document/8160759>
- [49] P. J. Lang and M. Bradley, "The int. affective picture system (iaps) in the study of emotion and attention," in *Handbook of Emotion Elicitation and Attention*, J. A. Coan and J. B. Allen, Eds. Oxford University Press, 2007, ch. 2, pp. 29–46.
- [50] A. Kushki, J. Fairley, S. Merja, G. King, and T. Chau, "Comparison of blood volume pulse and skin conductance responses to mental and affective stimuli at different anatomical sites," *Physiological measurement*, vol. 32, no. 10, p. 1529, 2011.
- [51] S. Greene, H. Thapliyal, and A. Caban-Holt, "A survey of affective computing for stress detection: Evaluating technologies in stress detection for better health," *IEEE Consumer Electronics Magazine*, vol. 5, no. 4, pp. 44–56, 2016.
- [52] M. B. Akçay and K. Oğuz, "Speech emotion recognition: Emotional models, databases, features, preprocessing methods, supporting modalities, and classifiers," *Speech Communication*, vol. 116, pp. 56–76, 2020.
- [53] B. W. Schuller, "Speech emotion recognition: Two decades in a nutshell, benchmarks, and ongoing trends," *Communications of the ACM*, vol. 61, no. 5, pp. 90–99, 2018.
- [54] K. Sailunaz, M. Dhaliwal, J. Rokne, and R. Alhajj, "Emotion detection from text and speech: a survey," *Social Network Analysis and Mining*, vol. 8, no. 1, p. 28, 2018.
- [55] D. Issa, M. F. Demirci, and A. Yazici, "Speech emotion recognition with deep convolutional neural networks," *Biomedical Signal Processing and Control*, vol. 59, p. 101894, 2020.
- [56] H. Beigi, "Speaker recognition," in *Fundamentals of Speaker Recognition*. Springer, 2011, pp. 543–559.
- [57] C. M. Lee and S. S. Narayanan, "Toward detecting emotions in spoken dialogs," *IEEE transactions on speech and audio processing*, vol. 13, no. 2, pp. 293–303, 2005.
- [58] D. Ververidis and C. Kotropoulos, "Emotional speech recognition: Resources, features, and methods," *Speech communication*, vol. 48, no. 9, pp. 1162–1181, 2006.
- [59] L. Chen, X. Mao, Y. Xue, and L. L. Cheng, "Speech emotion recognition: Features and classification models," *Digital signal processing*, vol. 22, no. 6, pp. 1154–1160, 2012.
- [60] B. Schuller, G. Rigoll, and M. Lang, "Speech emotion recognition combining acoustic features and linguistic information in a hybrid support vector machine-belief network architecture," in *2004 IEEE International Conference on Acoustics, Speech, and Signal Processing*, vol. 1. IEEE, 2004, pp. I–577.
- [61] Z. Zhao, Z. Bao, Y. Zhao, Z. Zhang, N. Cummins, Z. Ren, and B. Schuller, "Exploring deep spectrum representations via attention-based recurrent and convolutional neural networks for speech emotion recognition," *IEEE Access*, vol. 7, pp. 97 515–97 525, 2019.
- [62] G. Trigeorgis, F. Ringeval, R. Brueckner, E. Marchi, M. A. Nicolaou, B. Schuller, and S. Zafeiriou, "Adieu features? end-to-end speech emotion recognition using a deep convolutional recurrent network," in *2016 IEEE international conference on acoustics, speech and signal processing (ICASSP)*. IEEE, 2016, pp. 5200–5204.
- [63] S. Otl, S. Amiriparian, M. Gerczuk, V. Karas, and B. Schuller, "Group-level speech emotion recognition utilising deep spectrum features," in *Proceedings of the 2020 International Conference on Multimodal Interaction*, 2020, pp. 821–826.
- [64] L. Chao, J. Tao, M. Yang, Y. Li, and Z. Wen, "Long short term memory recurrent neural network based multimodal dimensional emotion recognition," in *Proceedings of the 5th International Workshop on Audio/Visual Emotion Challenge*, 2015, pp. 65–72.

- [65] M. Ren, W. Nie, A. Liu, and Y. Su, "Multi-modal correlated network for emotion recognition in speech," *Visual Informatics*, vol. 3, no. 3, pp. 150–155, 2019.
- [66] S. Ramakrishnan and I. M. El Emary, "Speech emotion recognition approaches in human computer interaction," *Telecommunication Systems*, vol. 52, no. 3, pp. 1467–1478, 2013.
- [67] A. Nanda, P. K. Sa, S. K. Choudhury, S. Bakshi, and B. Majhi, "A neuromorphic person re-identification framework for video surveillance," *IEEE Access*, vol. 5, pp. 6471–6482, 2017.
- [68] W. Han, T. Jiang, Y. Li, B. Schuller, and H. Ruan, "Ordinal learning for emotion recognition in customer service calls," in *ICASSP 2020-2020 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2020, pp. 6494–6498.
- [69] B. Li, D. Dimitriadis, and A. Stolcke, "Acoustic and lexical sentiment analysis for customer service calls," in *ICASSP 2019-2019 IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP)*. IEEE, 2019, pp. 5876–5880.
- [70] D. Morrison, R. Wang, and L. C. De Silva, "Ensemble methods for spoken emotion recognition in call-centres," *Speech communication*, vol. 49, no. 2, pp. 98–112, 2007.
- [71] A. Batliner, K. Fischer, J. Huber, J. Spilker, and E. Nöth, "How to find trouble in communication," *Speech communication*, vol. 40, no. 1-2, pp. 117–143, 2003.
- [72] V. Petrushin, "Emotion in speech: Recognition and application to call centers," in *Proceedings of artificial neural networks in engineering*, vol. 710, 1999, p. 22.
- [73] C. Parnin, "Subvocalization - toward hearing the inner thoughts of developers," 07 2011, pp. 197 – 200.
- [74] D. Fucci, D. Girardi, N. Novielli, L. Quaranta, and F. Lanubile, "A replication study on code comprehension and expertise using lightweight biometric sensors," in *Proc. of the 27th Int. Conf. on Program Comprehension, ICPC 2019, Montreal, QC, Canada, May 25-31, 2019*, 2019, pp. 311–322. [Online]. Available: <https://dl.acm.org/citation.cfm?id=3339126>
- [75] M. Züger, S. C. Müller, A. N. Meyer, and T. Fritz, "Sensing interruptibility in the office: A field study on the use of biometric and computer interaction sensors," in *Proc. of the 2018 CHI Conf. on Human Factors in Computing Systems, (CHI 2018)*, 2018, p. 591. [Online]. Available: <https://doi.org/10.1145/3173574.3174165>
- [76] S. C. Müller and T. Fritz, "Using (bio)metrics to predict code quality online," in *Proc. of the 38th Int. Conf. on Software Engineering, ICSE 2016, Austin, TX, USA, May 14-22, 2016*, 2016, pp. 452–463. [Online]. Available: <https://doi.org/10.1145/2884781.2884803>
- [77] S. A. Scherr, P. Mennig, C. Kammler, and F. Elberzhager, "On the road to enriching the app improvement process with emotions," in *2019 IEEE 27th Int. Requirements Engineering Conf. Workshops (REW)*. IEEE, 2019, pp. 84–91.
- [78] J. O. Johanssen, J. P. Bernius, and B. Bruegge, "Toward usability problem identification based on user emotions derived from facial expressions," in *2019 IEEE/ACM 4th Int. Workshop on Emotion Awareness in Software Engineering (SEmotion)*. IEEE, 2019, pp. 1–7.
- [79] M. Stade, S. A. Scherr, P. Mennig, F. Elberzhager, and N. Seyff, "Don't worry, be happy—exploring users' emotions during app usage for requirements engineering," in *2019 IEEE 27th Int. Requirements Engineering Conf. (RE)*. IEEE, 2019, pp. 375–380.
- [80] A. Sutcliffe, "Emotional requirements engineering," in *2011 IEEE 19th Int. Requirements Engineering Conf.* IEEE, 2011, pp. 321–322.
- [81] E. Guzman, R. Alkadhi, and N. Seyff, "An exploratory study of twitter messages about software applications," *Requirements Engineering*, vol. 22, no. 3, pp. 387–412, 2017.
- [82] D. Martens and W. Maalej, "Release early, release often, and watch your users' emotions: Lessons from emotional patterns," *IEEE Software*, vol. 36, no. 5, pp. 32–37, 2019.
- [83] C. Werner, Z. S. Li, and D. Damian, "Can a machine learn through customer sentiment?: A cost-aware approach to predict support ticket escalations," *IEEE Software*, vol. 36, no. 5, pp. 38–45, 2019.
- [84] R. Colomo-Palacios, C. Casado-Lumbreras, P. Soto-Acosta, and Á. García-Crespo, "Using the affect grid to measure emotions in software requirements engineering," *The International Journal of Universal Computer Science*, 2011.
- [85] T. Miller, S. Pedell, A. A. Lopez-Lorca, A. Mendoza, L. Sterling, and A. Keirnan, "Emotion-led modelling for people-oriented requirements engineering: the case study of emergency systems," *Journal of Systems and Software*, vol. 105, pp. 54–71, 2015.
- [86] P. Kamthan and N. Shahmir, "Effective user stories are affective," in *Int. Conf. on Ubiquitous Computing and Ambient Intelligence*. Springer, 2017, pp. 605–611.
- [87] B. Bent, B. A. Goldstein, W. A. Kibbe, and J. P. Dunn, "Investigating sources of inaccuracy in wearable optical heart rate sensors," *NPJ digital medicine*, vol. 3, no. 1, pp. 1–9, 2020.
- [88] H. Critchley and Y. Nagai, *Electrodermal Activity (EDA)*. New York, NY: Springer New York, 2013, pp. 666–669.
- [89] J. Allen, "Photoplethysmography and its application in clinical physiological measurement," *Physiological measurement*, vol. 28, no. 3, 2007.
- [90] J. E. Sinex, "Pulse oximetry: principles and limitations," *The American journal of emergency medicine*, vol. 17, no. 1, pp. 59–66, 1999.
- [91] B. Schuller and A. Batliner, *Computational Paralinguistics: Emotion, Affect and Personality in Speech and Language Processing*. Chichester: John Wiley & Sons., 2013.
- [92] E. S. Dan-Glauser and K. R. Scherer, "The geneva affective picture database (gaped): a new 730-picture database focusing on valence and normative significance," *Behavior research methods*, vol. 43, no. 2, p. 468, 2011.
- [93] J. J. Braithwaite, D. G. Watson, R. Jones, and M. Rowe, "A Guide for Analysing Electrodermal Activity (EDA) & Skin Conductance Responses (SCRs) for Psychological Experiments," University of Birmingham, UK, University of Birmingham, UK, Tech. Rep., 2015.
- [94] A. Greco, G. Valenza, A. Lanata, E. P. Scilingo, and L. Citi, "cvxEDA: A Convex Optimization Approach to Electrodermal Activity Processing," *IEEE Transactions on Biomedical Engineering*, vol. 63, no. 4, pp. 797–804, April 2016.
- [95] B. McFee, C. Raffel, D. Liang, D. P. Ellis, M. McVicar, E. Battenberg, and O. Nieto, "librosa: Audio and music signal analysis in python." Citeseer, 2015.
- [96] C. Tantithamthavorn, S. McIntosh, A. E. Hassan, and K. Matsumoto, "Automated parameter optimization of classification techniques for defect prediction models," in *Proc. of the 38th Int. Conf. on Software Engineering, ICSE 2016, Austin, TX, USA, May 14-22, 2016*, 2016, pp. 321–332.
- [97] —, "The impact of automated parameter optimization on defect prediction models," *IEEE Transactions on Software Engineering*, vol. 45, no. 7, pp. 683–711, July 2019.
- [98] S. Raschka, "Model evaluation, model selection, and algorithm selection in machine learning," *CoRR*, vol. abs/1811.12808, 2018. [Online]. Available: <http://arxiv.org/abs/1811.12808>
- [99] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: synthetic minority over-sampling technique," *Journal of artificial intelligence research*, vol. 16, pp. 321–357, 2002.
- [100] F. Sebastiani, "Machine learning in automated text categorization," *ACM Computing Surveys*, vol. 34, pp. 1–47, 04 2001.
- [101] M. Trice, L. Potts, and R. Small, "Values versus rules in social media communities: How platforms generate amorality on reddit and facebook," in *Digital Ethics*. Routledge, 2019, pp. 33–50.
- [102] P. Laukka, *Vocal Communication of Emotion*. Cham: Springer International Publishing, 2017, pp. 1–6. [Online]. Available: https://doi.org/10.1007/978-3-319-28099-8_562-1
- [103] K. R. Scherer, "Vocal communication of emotion: A review of research paradigms," *Speech Communication*, vol. 40, no. 1, pp. 227–256, 2003. [Online]. Available: <https://www.sciencedirect.com/science/article/pii/S0167639302000845>
- [104] M. Soleymani, M. Pantic, and T. Pun, "Multimodal emotion recognition in response to videos (extended abstract)," in *2015 Int. Conf. on Affective Computing and Intelligent Interaction, ACII 2015, Xi'an, China, September 21-24, 2015*, Sep. 2015, pp. 491–497.
- [105] M. Chen, J. Han, L. Guo, J. Wang, and I. Patras, "Identifying valence and arousal levels via connectivity between eeg channels," in *Proc. of the 2015 Int. Conf. on Affective Computing and Intelligent Interaction (ACII)*, ser. ACII '15. USA: IEEE Computer Society, 2015, pp. 63–69.
- [106] H. F. García, M. A. Álvarez, and Á. Á. Orozco, "Gaussian process dynamical models for multimodal affect recognition," in *38th Annual Int. Conf. of the IEEE Engineering in Medicine and Biology Society, EMBC 2016, Aug 2016*, pp. 850–853.
- [107] M. Aledhari, R. Razzak, R. M. Parizi, and G. Srivastava, "Deep neural networks for detecting real emotions using biofeedback and voice," in *Pattern Recognition. ICPR International Workshops and Challenges - Virtual Event, January 10-15, 2021, Proceedings, Part IV*, ser. Lecture Notes in Computer Science,

- A. D. Bimbo, R. Cucchiara, S. Sclaroff, G. M. Farinella, T. Mei, M. Bertini, H. J. Escalante, and R. Vezzani, Eds., vol. 12664. Springer, 2020, pp. 302–309. [Online]. Available: https://doi.org/10.1007/978-3-030-68799-1_21
- [108] A. Distanont, H. Haapasalo, M. Vaananen, and J. Lehto, “The engagement between knowledge transfer and requirements engineering,” *IJKL*, vol. 1, no. 2, pp. 131–156, 2012.
- [109] A. Haag, S. Goronzy, P. Schaich, and J. Williams, “Emotion recognition using bio-sensors: First steps towards an automatic system,” in *ADS’04*. Springer, 2004, pp. 36–48.
- [110] J. Wagner, J. Kim, and E. Andre, “From Physiological Signals to Emotions: Implementing and Comparing Selected Methods for Feature Extraction and Classification,” in *ICME 2005*. IEEE, 2005.
- [111] R. Mandryk, K. Inkpen, and T. Calvert, “Using psychophysiological techniques to measure user experience with entertainment technologies,” *Behaviour & Information Technology*, vol. 25, no. 2, pp. 141–158, 2006.
- [112] J. Scheirer, R. Fernandez, J. Klein, and R. Picard, “Frustrating the user on purpose: a step toward building an affective computer,” *Interacting with Computers*, Jan. 2002.
- [113] G. Zhou, J. H. Hansen, and J. F. Kaiser, “Nonlinear feature based classification of speech under stress,” *IEEE Transactions on speech and audio processing*, vol. 9, no. 3, pp. 201–216, 2001.
- [114] S. R. Bandela and T. K. Kumar, “Stressed speech emotion recognition using feature fusion of teager energy operator and mfcc,” in *2017 8th International Conference on Computing, Communication and Networking Technologies (ICCCNT)*. IEEE, 2017, pp. 1–5.
- [115] F. Zieris and L. Prechelt, “Two elements of pair programming skill,” in *2021 IEEE/ACM 43rd International Conference on Software Engineering: New Ideas and Emerging Results (ICSE-NIER)*, May 2021, pp. 51–55.
- [116] J. Kim and E. André, “Emotion recognition using physiological and speech signal in short-term observation,” *Lecture Notes in Computer Science (including subseries Lecture Notes in Artificial Intelligence and Lecture Notes in Bioinformatics)*, vol. 4021 LNAI, pp. 53–64, 2006.
- [117] A. Milne and N. Maiden, “Power and politics in requirements engineering: embracing the dark side?” *Requirements Engineering*, vol. 17, no. 2, pp. 83–98, 2012.
- [118] V. Gervasi, R. Gacitua, M. Rouncefield, P. Sawyer, L. Kof, L. Ma, P. Piwek, A. De Roeck, A. Willis, H. Yang *et al.*, “Unpacking tacit knowledge for requirements engineering,” in *Managing requirements knowledge*. Springer, 2013, pp. 23–47.
- [119] A. Ferrari, P. Spoletini, and S. Gnesi, “Ambiguity cues in requirements elicitation interviews,” in *2016 IEEE 24th International Requirements Engineering Conference (RE)*. IEEE, 2016, pp. 56–65.
- [120] A. Sutcliffe, P. Sawyer, G. Stringer, S. Couth, L. J. Brown, A. Gledson, C. Bull, P. Rayson, J. Keane, X.-j. Zeng *et al.*, “Known and unknown requirements in healthcare,” *Requirements engineering*, vol. 25, no. 1, pp. 1–20, 2020.
- [121] C. Wohlin, P. Runeson, M. Höst, M. C. Ohlsson, and B. Regnell, *Experimentation in software engineering*. Springer Science & Business Media, 2012.
- [122] J. G. Adair, “The hawthorne effect: a reconsideration of the methodological artifact,” *Journal of applied psychology*, vol. 69, no. 2, p. 334, 1984.
- [123] D. Girardi, F. Lanubile, N. Novielli, L. Quaranta, and A. Serebrenik, “Towards recognizing the emotions of developers using biometrics: The design of a field study,” in *Proc. of the 4th Int. Workshop on Emotion Awareness in Software Engineering, SEmotion@ICSE 2019, Montreal, QC, Canada, May 28, 2019*, 2019, pp. 13–16.
- [124] C. Tantithamthavorn, S. McIntosh, A. E. Hassan, and K. Matsumoto, “The impact of automated parameter optimization on defect prediction models,” *IEEE Transactions on Software Engineering*, vol. 45, no. 7, pp. 683–711, 2018.
- [125] M. Funk, “Cambridge analytica and the secret agenda of a facebook quiz,” *The New York Times*, 2016. [Online]. Available: <https://www.nytimes.com/2016/11/20/opinion/cambridge-analytica-facebook-quiz.html>
- [126] H. Davies, “Ted cruz campaign using firm that harvested data on millions of unwitting facebook users,” *The Guardian*. Archived, 2015. [Online]. Available: <https://www.theguardian.com/us-news/2015/dec/11/senator-ted-cruz-president-campaign-facebook-user-data>